

hasklib: Stochastic Process Derivations

Hubert Kasinski

August 25, 2026

Abstract

This document collects the stochastic-calculus results that the stochastic layer of **hasklib** relies on: from the probability-space setup, Brownian motion, Itô SDE, to Geometric Brownian Motion , Ornstein–Uhlenbeck, Arithmetic Brownian Motion, Cox–Ingersoll–Ross, and numerical schemes. Standard measure-theoretic probability is assumed; see ? or ? for construction.

Contents

Introduction

Thesis scope

This document is the mathematical companion to `hasklib`, a C++ library built around two intertwined purposes: a thesis spine on nonlinear filtering applied to commodity and interest-rate models (Schwartz–Smith, G2++), and a pricing showcase in which multiple engines — analytic, Monte Carlo, PDE, binomial tree — converge on the same vanilla option price. The chapters that follow build the stochastic-calculus foundations the filtering layer depends on: Itô calculus, then the four processes currently implemented (????, plus ABM and CIR), then the numerical schemes used to simulate them.

Library architecture

The `stochastic/` module is the only part of `hasklib` implemented and tested at the time of writing. It separates two concerns: *what* a process is (its drift and diffusion coefficients) from *how* it is simulated (the discretisation scheme). This mirrors the mathematical separation in ???? between the SDE itself and the numerical schemes of the final chapter.

Process hierarchy

Every process derives from an abstract `Process` base exposing only `drift(t,x)` and `diffusion(t,x)` — exactly the pair (a,b) of ???. GBM and ABM share a (μ, σ) parametrisation; OU and CIR share a (κ, θ, σ) mean-reverting one.

Scheme–process composition

Discretisation is deliberately *not* a virtual method on `Process`: `Scheme` holds a `Process` by reference rather than inheriting from it, so any scheme can drive any process without either class knowing about the other’s siblings. This is the software counterpart of the adapted-process argument underlying the Euler–Maruyama scheme, given in full in the final chapter.

Document conventions

Only results the library depends on downstream are stated in full; derivations not yet needed are left as **TO DERIVE** placeholders (see `TODO` in the source). The two diagrams above reflect the state of `stochastic/` only — `filter/`, `optim/`, and the `pybind11` layer are unbuilt and intentionally absent.




Figure 1: The **Process** hierarchy. Abstract operations (*italic*) are the interface consumed by both the exact samplers of **????** and the numerical schemes below.




Figure 2: Aggregation, not inheritance: a **Scheme** holds a **Process** reference rather than deriving from it, decoupling the two extension axes.

Chapter 1

Foundations: Itô Calculus and the SDE Framework

1.1 Probability space and Brownian motion

We work on a filtered probability space $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}_{t \geq 0}, \mathbb{P})$ satisfying the usual conditions. The filtration $\{\mathcal{F}_t\}$ represents the information available up to time t .

Definition 1.1 (Brownian motion).

A standard Brownian motion $\{W_t\}_{t \geq 0}$ is a process with

1. $W_0 = 0$ almost surely;
2. Independent increments: for $s < t < u$, $W_u - W_t$ is independent of $W_t - W_s$;
3. Gaussian increments: $W_t - W_s \sim \mathcal{N}(0, t - s)$;
4. Continuous sample paths.

Remark 1.1.

In code, property (3) is what `NormalRng::draw()` supplies: a step of length Δt uses $W_{t+\Delta t} - W_t = \sqrt{\Delta t} Z$ with $Z \sim \mathcal{N}(0, 1)$. This substitution is the single bridge between the continuous theory here and the discrete simulation code.

One might notice (2) and (4) pull in different directions. Independent increments mean each step of the path is a fresh draw, carrying no memory of the past — nothing about $\mathcal{F}_t$ constrains $W_{t+h} - W_t$. Continuity, on the other hand, means the path never jumps: W_{t+h} must stay close to W_t as $h \rightarrow 0$. It is not obvious a single process can do both at once, and the price of doing so turns out to be that the path is extremely rough and nowhere differentiable.

Fixing t and considering the difference quotient over a step $h > 0$. By the Gaussian-increment property, $W_{t+h} - W_t \sim \mathcal{N}(0, h)$, so the difference quotient

$$\frac{W_{t+h} - W_t}{h}$$

is itself Gaussian with mean 0 and variance

$$\text{Var}\left(\frac{W_{t+h} - W_t}{h}\right) = \frac{1}{h^2} \text{Var}(W_{t+h} - W_t) = \frac{h}{h^2} = \frac{1}{h},$$

using that dividing a random variable by a constant h scales its variance by $1/h^2$. Hence the standard deviation of the difference quotient is $1/\sqrt{h}$, which diverges as $h \downarrow 0$: the path moves

by order $\sqrt{h}$ over a step of order h , so no finite slope survives the limit. This informal argument is made rigorous by the Paley–Wiener–Zygmund theorem (?): with probability one, $t \mapsto W_t$ is nowhere differentiable. It also follows that the path has infinite first-order (total) variation on every interval, so integrals $\int f dW_s$ cannot be defined pathwise as Riemann–Stieltjes integrals. Its quadratic variation, however, is finite, and it is precisely this finite second-order variation that the Itô integral and Itô’s lemma are built on.

1.2 Quadratic variation

The property that distinguishes stochastic from ordinary calculus is that Brownian motion accumulates quadratic variation at a nonzero rate.

Proposition 1.1 (Quadratic variation of Brownian motion).

For a standard Brownian motion, the quadratic variation over $[0, T]$ satisfies

$$\langle W \rangle_T = \lim_{\|\Pi\| \rightarrow 0} \sum_k (W_{t_{k+1}} - W_{t_k})^2 = T \quad \text{in } L^2,$$

where Π is a partition of $[0, T]$. Informally this is written

$$(dW_t)^2 = dt.$$

Proof.

Write the partition as $\Pi = \{0 = t_0 < t_1 < \dots < t_n = T\}$ and abbreviate

$$\Delta t_k = t_{k+1} - t_k, \quad \Delta W_k = W_{t_{k+1}} - W_{t_k}, \quad S_\Pi = \sum_{k=0}^{n-1} (\Delta W_k)^2.$$

“Convergence in L^2 ” means the mean-square error vanishes: $\mathbb{E}[(S_\Pi - T)^2] \rightarrow 0$ as the mesh $\|\Pi\| := \max_k \Delta t_k \rightarrow 0$. We use only two facts about the increments, both directly from the definition of Brownian motion:

- Each $\Delta W_k \sim \mathcal{N}(0, \Delta t_k)$.
- Increments over disjoint intervals are independent.

We then calculate the mean of S_Π , i.e $\mathbb{E}[S_\Pi]$

Since $\mathbb{E}[\Delta W_k] = 0$, the second moment is the variance:

$$\mathbb{E}[(\Delta W_k)^2] = \text{Var}(\Delta W_k) = \Delta t_k.$$

Summing over k ,

$$\mathbb{E}[S_\Pi] = \sum_k \mathbb{E}[(\Delta W_k)^2] = \sum_k \Delta t_k = T.$$

This holds for every partition, not merely in the limit. Because the mean is exactly T , the mean-square error is the variance:

$$\mathbb{E}[(S_\Pi - T)^2] = \text{Var}(S_\Pi).$$

The whole statement then reduces to showing $\text{Var}(S_\Pi) \rightarrow 0$.

We then calculate the variance of a single term.

Normalising the increment as $\Delta W_k = \sqrt{\Delta t_k} Z$ with $Z \sim \mathcal{N}(0, 1)$, so $(\Delta W_k)^2 = \Delta t_k Z^2$ and

$$\text{Var}((\Delta W_k)^2) = (\Delta t_k)^2 \text{Var}(Z^2).$$

For a standard normal, $\mathbb{E}[Z^2] = 1$ and $\mathbb{E}[Z^4] = 3$, hence $\text{Var}(Z^2) = \mathbb{E}[Z^4] - (\mathbb{E}[Z^2])^2 = 3 - 1 = 2$. (The fourth moment is one Gaussian integration by parts: with the density ϕ satisfying $\phi'(z) = -z\phi(z)$, integrating $\int z^4\phi dz$ by parts gives $\mathbb{E}[Z^4] = 3\mathbb{E}[Z^2] = 3$.) Therefore,

$$\text{Var}((\Delta W_k)^2) = 2(\Delta t_k)^2.$$

Step 3 Variance of the sum, then let the mesh shrink.

The increments ΔW_k are independent, meaning the squared increments $(\Delta W_k)^2$ are independent too, and the variance of a sum of independent variables is the sum of the variances:

$$\text{Var}(S_\Pi) = \sum_k \text{Var}((\Delta W_k)^2) = 2 \sum_k (\Delta t_k)^2.$$

Bounding one factor of Δt_k in each term by the mesh:

$$\sum_k (\Delta t_k)^2 \leq \left(\max_k \Delta t_k \right) \sum_k \Delta t_k = \|\Pi\| T.$$

Combining,

$$\mathbb{E}[(S_\Pi - T)^2] = \text{Var}(S_\Pi) \leq 2 \|\Pi\| T \xrightarrow{\|\Pi\| \rightarrow 0} 0,$$

which is exactly the claimed L^2 convergence. $\square$

Remark 1.2.

The mean computation ($\mathbb{E}[S_\Pi] = T$) shows the average comes out exactly right; the variance computation shows it comes out reliably, not just on average. This second part matters because a single term is not close to its own mean: $(\Delta W_k)^2$ has standard deviation $\sqrt{2}\Delta t_k$, about the same size as its mean Δt_k itself. The reason the full sum S_Π still ends up close to T is that the terms are independent, so their ups and downs cancel out when added together — this cancellation is what the variance computation shows, and it is a property of the sum, not of any single term. Inside a sum (equivalently, an integral), the random $(\Delta W_k)^2$ may therefore be replaced by the deterministic Δt_k : this replacement is the entire content of the shorthand $(dW_t)^2 = dt$ used in the next section.

1.3 The Itô stochastic differential equation

Definition 1.2 (Itô SDE).

An Itô process X_t solves the stochastic differential equation

$$dX_t = a(t, X_t) dt + b(t, X_t) dW_t, \tag{1.1}$$

interpreted as the integral equation

$$X_t = X_0 + \int_0^t a(s, X_s) ds + \int_0^t b(s, X_s) dW_s,$$

where the second integral is an Itô integral. We call a the drift and b the diffusion coefficients.

Remark 1.3 (Correspondence with the code).

Equation (??) is exactly what the `StochasticProcess` base class encodes: `drift(t, x)` returns $a(t, x)$ and `diffusion(t, x)` returns $b(t, x)$. Every subsequent process (GBM, OU, CIR) is a choice of these two functions.

Assumption 1.1 (Existence and uniqueness).

We assume a and b satisfy the standard Lipschitz and linear-growth conditions guaranteeing a unique strong solution to (??); see ?, Thm. 5.2.1.

1.4 Itô's lemma

Itô's lemma is the chain rule for functions of an Itô process, and it is the engine behind every closed-form solution used in `hasklib` (the GBM exact sampler, the OU solution).

Theorem 1.1 (Itô's lemma, one dimension).
Let X_t solve (??) and let $\varphi(t, x) \in C^{1,2}$. Then

$$d\varphi(t, X_t) = \left(\frac{\partial \varphi}{\partial t} + a \frac{\partial \varphi}{\partial x} + \frac{1}{2} b^2 \frac{\partial^2 \varphi}{\partial x^2} \right) dt + b \frac{\partial \varphi}{\partial x} dW_t, \quad (1.2)$$

Informal derivation.

This is the standard second-order Taylor argument. It is not a proof in the strict sense, as a rigorous version requires the construction of the Itô integral (? , Sec. 4.2.1), which is outside the scope of this document. The derivation stems around the idea that, because dX carries a Brownian term of size $\sqrt{dt}$, a contribution that ordinary calculus discards as “second order” is in fact first order and should be kept.

Over a step of length dt the increment dW_t has typical size $\sqrt{dt}$ (it is $\mathcal{N}(0, dt)$).

Ranking the elementary products by their power of dt :

$$dt = O(dt), \quad dW_t = O(dt^{1/2}), \quad (dW_t)^2 = O(dt), \quad dt dW_t = O(dt^{3/2}), \quad (dt)^2 = O(dt^2).$$

Discarding everything of order smaller than dt , and using ?? to replace $(dW_t)^2$ by its in-sum limit dt , gives the Itô multiplication table

$$(dW_t)^2 = dt, \quad dt dW_t = 0, \quad (dt)^2 = 0. \quad (1.3)$$

The last two entries are not literally zero; they are negligible next to dt and vanish in the mesh limit for exactly the reason the fluctuations did in ??.

Taylor expansion:

Expand $\varphi(t + dt, X_t + dX)$ about (t, X_t) to second order:

$$d\varphi = \frac{\partial \varphi}{\partial t} dt + \frac{\partial \varphi}{\partial x} dX + \frac{1}{2} \frac{\partial^2 \varphi}{\partial t^2} (dt)^2 + \frac{\partial^2 \varphi}{\partial t \partial x} dt dX + \frac{1}{2} \frac{\partial^2 \varphi}{\partial x^2} (dX)^2 + \dots$$

Substitute the SDE $dX = a dt + b dW_t$ and square:

$$(dX)^2 = a^2 (dt)^2 + 2ab dt dW_t + b^2 (dW_t)^2.$$

Apply (??). In $(dX)^2$ the first two terms die and only $b^2 (dW_t)^2 = b^2 dt$ survives:

$$(dX)^2 = b^2 dt.$$

The two remaining second-order terms of the expansion vanish for the same reason:

$$\frac{1}{2} \frac{\partial^2 \varphi}{\partial t^2} (dt)^2 = 0, \quad \frac{\partial^2 \varphi}{\partial t \partial x} dt dX = \frac{\partial^2 \varphi}{\partial t \partial x} (a (dt)^2 + b dt dW_t) = 0.$$

Leaving:

$$d\varphi = \frac{\partial \varphi}{\partial t} dt + \frac{\partial \varphi}{\partial x} dX + \frac{1}{2} \frac{\partial^2 \varphi}{\partial x^2} b^2 dt.$$

Substitute $dX = a dt + b dW_t$ once more and collect the dt and dW_t parts:

$$d\varphi = \left(\frac{\partial \varphi}{\partial t} + a \frac{\partial \varphi}{\partial x} + \frac{1}{2} b^2 \frac{\partial^2 \varphi}{\partial x^2} \right) dt + b \frac{\partial \varphi}{\partial x} dW_t,$$

which is (??). □

Remark 1.4.

Had X been a smooth (finite-variation) path, $(dX)^2$ would have been genuinely $O((dt)^2)$ and dropped out, leaving the ordinary chain rule $d\varphi = \frac{\partial\varphi}{\partial t} dt + \frac{\partial\varphi}{\partial x} dX$. The whole gap between stochastic and ordinary calculus is the single surviving term $\frac{1}{2}b^2\frac{\partial^2\varphi}{\partial x^2} dt$, and it survives only because $(dW_t)^2 = dt$ rather than 0 — i.e. because Brownian motion has nonzero quadratic variation. Every closed-form result in the code that “corrects a drift by half a variance” is exactly this one term.

Chapter 2

Geometric Brownian Motion

2.1 The GBM SDE

Definition 2.1 (Geometric Brownian motion).

X_t follows geometric Brownian motion with drift $\mu \in \mathbb{R}$ and volatility $\sigma > 0$ if it solves

$$dX_t = \mu X_t dt + \sigma X_t dW_t, \quad X_0 > 0. \quad (2.1)$$

Referring back to the previous chapter on Itô Calculus, this is the choice $a(t, x) = \mu x$, $b(t, x) = \sigma x$ in the general SDE $dX_t = a(t, X_t)dt + b(t, X_t)dW_t$.

Remark 2.1 (Correspondence with the code).

This is exactly why the pair of functions `GBM::drift` and `GBM::diffusion` return: `drift(t, x) =  $\mu x$` , `diffusion(t, x) =  $\sigma x$` . `EulerMaruyama` consumes these polymorphically through the `StochasticProcess` interface; the sampler derived below is the exact alternative to that discretisation, used by `GBM::sample_terminal` instead of stepping the Euler scheme forward.

Both a and b are affine in x , hence globally Lipschitz and of linear growth on all of $\mathbb{R}$ (??), so (??) has a unique strong solution for any $X_0 \in \mathbb{R}$. In particular this holds for $X_0 > 0$; positivity of X_t itself is established below once the explicit solution is derived.

2.2 Linearising via Itô's lemma

The SDE (??) has no constant-coefficient closed form as written, because both coefficients scale with X_t . The standard fix is to apply Itô's lemma to $\varphi(t, x) = \ln x$: this removes the x -dependence from both coefficients and turns (??) into an SDE with constant drift and diffusion, which integrates directly.

Theorem 2.1 (Exact terminal law of GBM).

Let X_t solve (??). Then for any $T > 0$,

$$X_T = X_0 \exp\left(\left(\mu - \frac{1}{2}\sigma^2\right)T + \sigma W_T\right), \quad (2.2)$$

where $W_T \sim \mathcal{N}(0, T)$.

Proof.

Apply Itô's lemma (??) to $\varphi(t, x) = \ln x$, so:

$\frac{\partial \varphi}{\partial t} = 0$, $\frac{\partial \varphi}{\partial x} = \frac{1}{x}$, $\frac{\partial^2 \varphi}{\partial x^2} = -\frac{1}{x^2}$. With $a(t, X_t) = \mu X_t$ and $b(t, X_t) = \sigma X_t$,

$$d(\ln X_t) = \left(0 + \mu X_t \cdot \frac{1}{X_t} + \frac{1}{2} \sigma^2 X_t^2 \cdot \left(-\frac{1}{X_t^2}\right)\right) dt + \sigma X_t \cdot \frac{1}{X_t} dW_t.$$

The X_t factors cancel completely, leaving

$$d(\ln X_t) = \left(\mu - \frac{1}{2}\sigma^2\right)dt + \sigma dW_t. \quad (2.3)$$

The right-hand side no longer depends on X_t : it is Brownian motion with constant drift $\mu - \frac{1}{2}\sigma^2$ and constant volatility σ , so (??) integrates directly from 0 to T :

$$\ln X_T - \ln X_0 = \left(\mu - \frac{1}{2}\sigma^2\right)T + \sigma(W_T - W_0) = \left(\mu - \frac{1}{2}\sigma^2\right)T + \sigma W_T,$$

using $W_0 = 0$. Exponentiating gives (??). Since $W_T \sim \mathcal{N}(0, T)$ by definition of Brownian motion, this is a complete description of the law of X_T given X_0 . $\square$

Remark 2.2.

Equation (??) is sampled with a single normal draw: generate $Z \sim \mathcal{N}(0, 1)$, set $W_T = \sqrt{T} Z$, and evaluate the right-hand side. No time-stepping is needed, which is what makes it an exact sampler rather than a discretisation, and it is exactly what `GBM::sample_terminal` implements using `NormalRng::draw()`.

Remark 2.3 (Where the $-\frac{1}{2}\sigma^2$ comes from).

The drift that appears in (??) is not μ , the drift of X_t itself, but $\mu - \frac{1}{2}\sigma^2$. That correction is exactly the $\frac{1}{2}b^2\frac{\partial^2\varphi}{\partial x^2}$ term of Itô's lemma, which is the term with no analogue in ordinary calculus, mentioned in the previous chapter. If one (incorrectly) applied the ordinary chain rule to $\ln X_t$, the correction would be missing and the resulting sampler would be biased upward in expectation, as ?? below makes precise.

2.3 Moments

The moment-matching test in `test_stochastic.cpp` checks simulated terminal samples against the following closed-form mean and variance.

Lemma 2.1 (Moment generating function of a Gaussian).

If $Y \sim \mathcal{N}(m, v)$ then $\mathbb{E}[e^Y] = \exp(m + \frac{1}{2}v)$.

Proof.

Write $Y = m + \sqrt{v} Z$ with $Z \sim \mathcal{N}(0, 1)$. Complete the square in the Gaussian integral:

$$\mathbb{E}[e^Y] = e^m \mathbb{E}[e^{\sqrt{v} Z}] = e^m \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}} e^{\sqrt{v} z - \frac{1}{2}z^2} dz = e^m e^{\frac{1}{2}v} \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(z - \sqrt{v})^2} dz = e^{m + \frac{1}{2}v},$$

the last integral being 1 since it is a $\mathcal{N}(\sqrt{v}, 1)$ density integrated over $\mathbb{R}$. $\square$

Proposition 2.1 (Mean of GBM).

$$\mathbb{E}[X_T] = X_0 e^{\mu T}.$$

Proof.

By ??, $X_T = X_0 e^Y$ with $Y = (\mu - \frac{1}{2}\sigma^2)T + \sigma W_T$.

Since $W_T \sim \mathcal{N}(0, T)$, $Y \sim \mathcal{N}((\mu - \frac{1}{2}\sigma^2)T, \sigma^2 T)$.

Applying ?? with $m = (\mu - \frac{1}{2}\sigma^2)T$ and $v = \sigma^2 T$,

$$\mathbb{E}[X_T] = X_0 \mathbb{E}[e^Y] = X_0 \exp\left((\mu - \frac{1}{2}\sigma^2)T + \frac{1}{2}\sigma^2 T\right) = X_0 e^{\mu T}.$$

Where the $-\frac{1}{2}\sigma^2 T$ Itô correction and the $+\frac{1}{2}\sigma^2 T$ from the moment generating function cancel exactly, leaving the mean growing at the "natural" rate μ despite the correction appearing inside every individual path. $\square$

Proposition 2.2 (Variance of GBM).

$$\text{Var}(X_T) = X_0^2 e^{2\mu T} (e^{\sigma^2 T} - 1).$$

Proof.

$$\text{Var}(X_T) = \mathbb{E}[X_T^2] - (\mathbb{E}[X_T])^2.$$

For the second moment, write $X_T^2 = X_0^2 e^{2Y}$ with Y as above, so $2Y \sim \mathcal{N}(2(\mu - \frac{1}{2}\sigma^2)T, 4\sigma^2 T)$. By ?? with $m = 2(\mu - \frac{1}{2}\sigma^2)T$ and $v = 4\sigma^2 T$,

$$\mathbb{E}[X_T^2] = X_0^2 \exp\left(2\left(\mu - \frac{1}{2}\sigma^2\right)T + 2\sigma^2 T\right) = X_0^2 e^{(2\mu + \sigma^2)T}.$$

Subtracting $(\mathbb{E}[X_T])^2 = X_0^2 e^{2\mu T}$ from ?? and factoring,

$$\text{Var}(X_T) = X_0^2 e^{(2\mu + \sigma^2)T} - X_0^2 e^{2\mu T} = X_0^2 e^{2\mu T} (e^{\sigma^2 T} - 1). \quad \square$$

Remark 2.4 (What the test checks).

In the code, `test_stochastic.cpp` draws a large number of terminal samples from `GBM::sample_terminal`, computes their sample mean and sample variance, and checks these against `????` to within Monte Carlo tolerance. This is a check on the sampler, not on `(?)` itself — it would catch, for instance, a sign error in the $-\frac{1}{2}\sigma^2$ correction, since that error changes the mean but leaves the qualitative shape of the distribution intact.

Chapter 3

The Ornstein–Uhlenbeck Process

3.1 The OU SDE

Definition 3.1 (Ornstein–Uhlenbeck process).

X_t follows an Ornstein–Uhlenbeck process with mean-reversion speed $\kappa > 0$, long-run mean $\theta \in \mathbb{R}$, and volatility $\sigma > 0$ if it solves

$$dX_t = \kappa(\theta - X_t) dt + \sigma dW_t, \quad X_0 \in \mathbb{R}. \quad (3.1)$$

In the notation of Itô from the first chapter, this is the choice $a(t, x) = \kappa(\theta - x)$, $b(t, x) = \sigma$ (constant) in the general SDE $dX_t = a(t, X_t)dt + b(t, X_t)dW_t$.

Remark 3.1 (Correspondence with the code).

`OU::drift` returns $\kappa(\theta - x)$ and `OU::diffusion` returns the constant σ , independent of x — unlike GBM, the diffusion here does not scale with the state. `EulerMaruyama` can still step this process through the same `StochasticProcess` interface, but as with GBM there is an exact alternative, derived below, that `OU::sample_terminal` uses instead.

Both coefficients are affine (in particular globally Lipschitz and of linear growth) in x , so the existence–uniqueness assumption of Itô’s theory applies unconditionally here — unlike GBM, there is no positivity constraint to track, since OU is free to take any real value.

3.2 Exact solution: the integrating-factor trick

Unlike GBM, taking a logarithm does nothing useful here since (??) is already linear in X_t , not multiplicative. The relevant substitution instead is the same one used for linear ODEs: multiply by the integrating factor $e^{\kappa t}$ and let Itô’s lemma absorb the product rule.

Theorem 3.1 (Exact terminal law of OU).

Let X_t solve (??). Then for any $T > 0$,

$$X_T = X_0 e^{-\kappa T} + \theta(1 - e^{-\kappa T}) + \sigma e^{-\kappa T} \int_0^T e^{\kappa s} dW_s. \quad (3.2)$$

Conditional on X_0 , X_T is Gaussian.

Proof.

Apply Itô’s lemma (??) to $\varphi(t, x) = e^{\kappa t}x$, so $\varphi_t = \kappa e^{\kappa t}x$, $\varphi_x = e^{\kappa t}$, $\varphi_{xx} = 0$. With $a(t, X_t) = \kappa(\theta - X_t)$ and $b(t, X_t) = \sigma$,

$$d(e^{\kappa t} X_t) = \left(\kappa e^{\kappa t} X_t + \kappa(\theta - X_t) e^{\kappa t} + 0 \right) dt + \sigma e^{\kappa t} dW_t.$$

The X_t terms cancel completely — $\kappa e^{\kappa t} X_t$ against $-\kappa X_t e^{\kappa t}$ — which is precisely the point of the integrating factor, leaving

$$d(e^{\kappa t} X_t) = \kappa \theta e^{\kappa t} dt + \sigma e^{\kappa t} dW_t. \quad (3.3)$$

The right-hand side no longer involves X_t at all, so (??) integrates directly from 0 to T :

$$e^{\kappa T} X_T - X_0 = \kappa \theta \int_0^T e^{\kappa s} ds + \sigma \int_0^T e^{\kappa s} dW_s = \theta(e^{\kappa T} - 1) + \sigma \int_0^T e^{\kappa s} dW_s.$$

Multiplying by $e^{-\kappa T}$ gives (??).

The stochastic integral $\int_0^T e^{\kappa s} dW_s$ is a Wiener integral of a deterministic integrand, hence Gaussian (a linear functional of a Gaussian process); X_T is an affine transformation of it and of the deterministic terms, so X_T is Gaussian as well. $\square$

Lemma 3.1 (Moments of a Wiener integral).

For deterministic $g \in L^2[0, T]$, the stochastic integral $I = \int_0^T g(s) dW_s$ satisfies $\mathbb{E}[I] = 0$ and $\text{Var}(I) = \int_0^T g(s)^2 ds$ (the Itô isometry).

Citation:

This is a standard property of the Itô integral, following from its construction as an L^2 -limit of Riemann-type sums of independent Gaussian increments; see ? or ?. (Cited, not proved, in the same spirit as the existence–uniqueness ??.) $\square$

3.3 Moments and the stationary law

Proposition 3.1 (Mean and variance of OU).

$$\mathbb{E}[X_T | X_0] = X_0 e^{-\kappa T} + \theta(1 - e^{-\kappa T}), \quad \text{Var}(X_T | X_0) = \frac{\sigma^2}{2\kappa} (1 - e^{-2\kappa T}).$$

Proof.

Take expectations in (??).

By ?? with $g(s) = \sigma e^{-\kappa T} e^{\kappa s}$ (constant $e^{-\kappa T}$ pulled outside the integral),

$\mathbb{E}[\int_0^T e^{\kappa s} dW_s] = 0$, so the mean is just the deterministic part:

$$\mathbb{E}[X_T | X_0] = X_0 e^{-\kappa T} + \theta(1 - e^{-\kappa T}).$$

For the variance, write $X_T = a + bI$ where $a = X_0 e^{-\kappa T} + \theta(1 - e^{-\kappa T})$ is deterministic given X_0 , $b = \sigma e^{-\kappa T}$, and $I = \int_0^T e^{\kappa s} dW_s$. Since $\text{Var}(a + bY) = b^2 \text{Var}(Y)$ for constants a, b ,

$$\text{Var}(X_T | X_0) = \sigma^2 e^{-2\kappa T} \text{Var}\left(\int_0^T e^{\kappa s} dW_s\right) = \sigma^2 e^{-2\kappa T} \int_0^T e^{2\kappa s} ds,$$

using ?? with $g(s) = e^{\kappa s}$. The integral evaluates to $(e^{2\kappa T} - 1)/(2\kappa)$, giving

$$\text{Var}(X_T | X_0) = \sigma^2 e^{-2\kappa T} \cdot \frac{e^{2\kappa T} - 1}{2\kappa} = \frac{\sigma^2}{2\kappa} (1 - e^{-2\kappa T}). \quad \square$$

Remark 3.2.

Both the mean and the variance are exactly what `OU::sample_terminal` needs: draw $Z \sim \mathcal{N}(0, 1)$ once and set

$$X_T = X_0 e^{-\kappa T} + \theta(1 - e^{-\kappa T}) + \sqrt{\frac{\sigma^2}{2\kappa} (1 - e^{-2\kappa T})} Z,$$

exactly as GBM's exact sampler needed only Z and the closed-form mean and variance from ?? and ??. No time-stepping, hence no discretisation bias.

Corollary 3.1 (Stationary law).

As $T \rightarrow \infty$ (with $\kappa > 0$ fixed), $X_T \mid X_0$ converges in distribution to

$$X_\infty \sim \mathcal{N}\left(\theta, \frac{\sigma^2}{2\kappa}\right),$$

independent of the starting point X_0 .

Proof.

Immediate from ?? : $e^{-\kappa T} \rightarrow 0$ and $e^{-2\kappa T} \rightarrow 0$ as $T \rightarrow \infty$ since $\kappa > 0$, so the mean $\rightarrow \theta$ and the variance $\rightarrow \sigma^2/(2\kappa)$, independent of X_0 . $\square$

Remark 3.3 (Why $\kappa > 0$ matters).

The mean-reversion speed $\kappa > 0$ is what makes $e^{-\kappa T} \rightarrow 0$ and hence guarantees a stationary law at all; it is also what makes OU a sensible model for a quantity that fluctuates around a long-run level θ rather than drifting away from it, in contrast to GBM's persistent (possibly explosive) drift. If $\kappa \leq 0$ the process is non-stationary and ?? does not apply — the library assumes $\kappa > 0$ throughout.

Remark 3.4 (Relevance beyond the pricing showcase).

OU is not only one of the four processes in the convergence demonstration; it is also the driving process behind the Hull–White short-rate model, $dr_t = \kappa(\theta(t) - r_t)dt + \sigma dW_t$, used throughout ? for term-structure modelling. The only difference from (??) is a time-dependent long-run level $\theta(t)$ calibrated to the observed yield curve, chosen precisely so that the affine structure exploited in ?? survives with θ replaced by $\theta(t)$ inside the integral.

Chapter 4

Arithmetic Brownian Motion

4.1 The ABM SDE

Definition 4.1 (Arithmetic Brownian motion).

X_t follows arithmetic Brownian motion with drift $\mu \in \mathbb{R}$ and volatility $\sigma > 0$ if it solves

$$dX_t = \mu dt + \sigma dW_t, \quad X_0 \in \mathbb{R}. \quad (4.1)$$

In the notation of Itô, this is $a(t, x) = \mu$, $b(t, x) = \sigma$ — both constant, with no x -dependence.

Remark 4.1 (Correspondence with the code).

`ABM::drift` returns the constant μ and `ABM::diffusion` returns the constant σ , independent of x in both cases. Since neither coefficient depends on x , the Lipschitz and linear-growth conditions from ??, hold trivially — this is the simplest possible case of that assumption.

4.2 Exact solution

Both GBM and OU needed a substitution to remove an X_t -dependence from the coefficients before the SDE could be integrated: $\ln x$ for GBM, the integrating factor $e^{\kappa t}$ for OU. Here there is nothing to remove — μ and σ are already constants — so the general relation from (??),

$$X_t = X_0 + \int_0^t a(s, X_s) ds + \int_0^t b(s, X_s) dW_s,$$

integrates immediately once $a = \mu$ and $b = \sigma$ are substituted.

Theorem 4.1 (Exact terminal law of ABM).

Let X_t solve (??). Then for any $T > 0$,

$$X_T = X_0 + \mu T + \sigma W_T, \quad (4.2)$$

where $W_T \sim \mathcal{N}(0, T)$, so

$$X_T \sim \mathcal{N}(X_0 + \mu T, \sigma^2 T). \quad (4.3)$$

Proof.

Substituting $a(s, X_s) = \mu$ and $b(s, X_s) = \sigma$,

$$X_T = X_0 + \int_0^T \mu ds + \int_0^T \sigma dW_s = X_0 + \mu T + \sigma \int_0^T dW_s.$$

The remaining integral is $\int_0^T dW_s = W_T - W_0 = W_T$, giving (??). Since $X_0 + \mu T$ is a fixed number and $W_T \sim \mathcal{N}(0, T)$, the affine transformation $X_0 + \mu T + \sigma W_T$ is itself Gaussian with mean $X_0 + \mu T$ and variance $\sigma^2 \text{Var}(W_T) = \sigma^2 T$, giving (??). $\square$

4.3 Moments and the relation to OU

?? already gives the moments directly, with no separate lemma needed:

$$\mathbb{E}[X_T] = X_0 + \mu T, \quad \text{Var}(X_T) = \sigma^2 T.$$

Unlike OU, the variance grows without bound as $T \rightarrow \infty$: there is no mean-reverting pull holding the process in, so uncertainty simply accumulates. Correspondingly, $\mathbb{E}[X_T]$ grows in a straight line rather than curving toward a long-run level — there is no θ here for it to curve toward.

Chapter 5

The Cox–Ingersoll–Ross Process

5.1 The CIR SDE

Definition 5.1 (Cox–Ingersoll–Ross process).

X_t follows a CIR process with mean-reversion speed $\kappa > 0$, long-run mean $\theta > 0$, and volatility $\sigma > 0$ if it solves

$$dX_t = \kappa(\theta - X_t) dt + \sigma\sqrt{X_t} dW_t, \quad X_0 > 0. \quad (5.1)$$

In the notation of Itô from the first chapter, this is $a(t, x) = \kappa(\theta - x)$ (identical to OU) and $b(t, x) = \sigma\sqrt{x}$.

Remark 5.1 (Correspondence with the code).

`CIR::drift` returns $\kappa(\theta - x)$, exactly matching `OU::drift`. `CIR::diffusion` returns $\sigma\sqrt{x}$, which requires $x \geq 0$ to be well-defined — unlike GBM’s σx or OU’s constant σ , this diffusion coefficient is not defined for negative arguments at all, which is the source of every complication in this file.

Remark 5.2 (A genuine gap in the existence–uniqueness story).

The existence–uniqueness ?? requires a and b to be Lipschitz. $\sqrt{x}$ is not Lipschitz near $x = 0$ (its slope $\rightarrow \infty$ there), so that assumption does not literally cover (??) as stated. A strong solution nonetheless exists and is unique — this requires a separate, more delicate existence–uniqueness theory for Hölder-continuous (rather than Lipschitz) diffusion coefficients; see ? or ?.

5.2 Why the earlier tricks do not apply here

GBM’s diffusion σx is proportional to the state, which $\ln x$ removed. OU’s diffusion σ is constant, which the integrating factor $e^{\kappa t}$ exploited alongside an affine drift. CIR’s drift $\kappa(\theta - x)$ is the same affine form as OU’s, so the integrating factor idea is tempting to reuse — but it was the diffusion term, not the drift, that made OU’s trick work cleanly: applying Itô’s lemma to $\varphi(t, x) = e^{\kappa t}x$ needs $\varphi_{xx} = 0$ to kill the second-order term entirely, which held because φ was linear in x . The moment we try to fold $\sigma\sqrt{x}$ into any single substitution $\varphi(t, x)$, the resulting φ_{xx} term no longer vanishes, and no elementary choice of φ linearises (??) the way $\ln x$ or $e^{\kappa t}x$ did for the earlier processes. This is not a failure of effort — $\sqrt{x}$ is a genuinely different kind of nonlinearity (a square-root, or CEV-type, diffusion), and no algebraic trick of this kind exists for it in general.

Remark 5.3 (The exact law exists, but is out of scope here).

CIR’s transition density is nonetheless known exactly: conditional on X_t , the (scaled) law of $X_{t+\Delta t}$ is a noncentral chi-squared distribution, a result following from CIR’s connection to squared Bessel processes (?). This is the route `CIR::sample_terminal` would take for an exact

sampler (Boost.Math's noncentral chi-squared machinery, per the library's `stats/` layer), analogous to `GBM::sample_terminal` and `OU::sample_terminal`. Deriving that transition density, however, requires either the squared-Bessel-process identification or a PDE/Laplace-transform argument neither established nor needed elsewhere in this document, so it is omitted here.

5.3 Moments, via an ODE rather than a path solution

Although no closed-form X_t is derived, its first two moments are obtainable directly: take expectations of the SDE (or of Itô's lemma applied to a power of X_t) and solve the resulting deterministic ODE. This sidesteps needing a path solution at all.

Proposition 5.1 (Mean of CIR).

Let $m(t) = \mathbb{E}[X_t | X_0]$. Then $m(t)$ solves $m'(t) = \kappa(\theta - m(t))$ with $m(0) = X_0$, giving

$$m(t) = X_0 e^{-\kappa t} + \theta(1 - e^{-\kappa t}),$$

identical in form to OU's mean (??).

Proof.

Taking expectations of the integral form of (??),

$$m(t) = X_0 + \kappa\theta \int_0^t ds - \kappa \int_0^t m(s) ds + \mathbb{E} \left[\int_0^t \sigma \sqrt{X_s} dW_s \right].$$

The stochastic integral has zero expectation (?? gives $\mathbb{E}[I] = 0$ for the underlying Wiener integral), regardless of the specific form of the diffusion coefficient inside it. Differentiating the remaining deterministic relation gives $m'(t) = \kappa\theta - \kappa m(t)$ with $m(0) = X_0$, exactly the deterministic ODE behind OU's mean (??), with the same solution:

$$\mathbb{E}[X_t | X_0] = X_0 e^{-\kappa t} + \theta(1 - e^{-\kappa t}). \quad \square$$

Remark 5.4.

The mean depends only on the drift, and CIR's drift is literally the same function as OU's; the two processes' means therefore coincide exactly, even though their diffusion coefficients — and hence their variances, paths, and positivity properties — are entirely different.

Proposition 5.2 (Variance of CIR).

Let $v(t) = \text{Var}(X_t | X_0)$. Then

$$v(t) = \frac{\sigma^2 X_0}{\kappa} (e^{-\kappa t} - e^{-2\kappa t}) + \frac{\sigma^2 \theta}{2\kappa} (1 - e^{-\kappa t})^2.$$

Proof. Apply Itô's lemma to $f(x) = x^2$: $f_x = 2x$, $f_{xx} = 2$, $f_t = 0$. With $a = \kappa(\theta - X_t)$ and $b = \sigma\sqrt{X_t}$,

$$d(X_t^2) = \left(2X_t \cdot \kappa(\theta - X_t) + \frac{1}{2} \cdot \sigma^2 X_t \cdot 2 \right) dt + 2X_t \cdot \sigma \sqrt{X_t} dW_t = (2\kappa\theta X_t - 2\kappa X_t^2 + \sigma^2 X_t) dt + 2\sigma X_t^{3/2} dW_t.$$

Taking expectations and writing $w(t) = \mathbb{E}[X_t^2]$, the stochastic term again vanishes in expectation, giving

$$w'(t) = (2\kappa\theta + \sigma^2)m(t) - 2\kappa w(t).$$

Rather than solve for $w(t)$ directly, substitute $w(t) = v(t) + m(t)^2$ and use $m'(t) = \kappa(\theta - m(t))$ from ?? to simplify. Differentiating $w = v + m^2$ gives $w' = v' + 2mm' = v' + 2\kappa\theta m - 2\kappa m^2$. Substituting into the equation for $w'(t)$ and cancelling the $2\kappa\theta m(t)$ terms on both sides,

$$v'(t) + 2\kappa\theta m(t) - 2\kappa m(t)^2 + 2\kappa(v(t) + m(t)^2) = (2\kappa\theta + \sigma^2)m(t) \implies v'(t) + 2\kappa v(t) = \sigma^2 m(t),$$

with $v(0) = 0$ (since X_0 is a fixed starting point, not itself random). This is a much cleaner ODE than the one for $w(t)$: the θ -dependence has been absorbed entirely into $m(t)$, which is already known. Solve by the same integrating-factor technique used for OU's SDE, now applied to this ordinary (non-stochastic) ODE: multiplying by $e^{2\kappa t}$ collapses the left side into $\frac{d}{dt}(e^{2\kappa t}v(t))$, and substituting $m(t)$ from ?? and integrating from 0 to t (using $v(0) = 0$) gives, after simplification,

$$v(t) = \frac{\sigma^2 X_0}{\kappa} (e^{-\kappa t} - e^{-2\kappa t}) + \frac{\sigma^2 \theta}{2\kappa} (1 - e^{-\kappa t})^2,$$

which is the claimed formula. (*The intermediate integrating-factor algebra is routine and omitted; it is the same mechanical steps as the OU variance derivation in ??, applied to this ODE instead of directly to the SDE.*) $\square$

Remark 5.5 (What this moment computation buys, and what it does not). ???? give exact first and second moments with no approximation — useful for moment-matching tests analogous to `test_stochastic.cpp`'s GBM check. They do not give the full distribution of X_t : CIR is not Gaussian (unlike OU or ABM), since $\sqrt{X_t}$ scales the noise unevenly across the state space. Recovering the full law requires the noncentral chi-squared transition density flagged above, not this ODE method.

5.4 The Feller condition

Theorem 5.1 (Feller condition). *If $2\kappa\theta \geq \sigma^2$, then $X_t > 0$ for all $t \geq 0$, almost surely, given $X_0 > 0$. If $2\kappa\theta < \sigma^2$, X_t can reach 0 in finite time (though it is instantaneously reflected back into $(0, \infty)$ by the drift, since $a(t, 0) = \kappa\theta > 0$).*

Citation. This is Feller's boundary classification for one-dimensional diffusions, applied to (??); see ? or ?. (*Cited, not proved — the proof requires the scale-function/speed-measure machinery for one-dimensional diffusions, which nothing else in this note depends on.*) $\square$

Remark 5.6 (Why this condition matters here specifically). Every other assumption in this companion (Lipschitz coefficients, $\kappa > 0$ for OU's stationary law) has been a technical convenience. This one is load-bearing in a different way: (??) is not even well-defined pathwise if X_t goes negative, since $\sqrt{X_t}$ has no real value there. The Feller condition is what guarantees the SDE never has to confront that question in continuous time. Nothing in `hasklib` currently enforces $2\kappa\theta \geq \sigma^2$ on constructing a CIR instance; whether to validate it, and what to do if a calibrated parameter set violates it (a common occurrence in practice), is a live design question, not a settled one.

Remark 5.7 (Relevance beyond the pricing showcase). CIR is, alongside Hull–White, one of the classical short-rate models in ? — valued for keeping rates strictly positive when the Feller condition holds, in contrast to Hull–White's Gaussian (and hence potentially negative) short rate. The parameter tension is real: market-calibrated CIR parameters frequently fail the Feller condition, which is precisely why the numerical scheme discussed next is not a mere implementation detail.

5.5 Motivation for full truncation

Even when the Feller condition holds and the true process X_t never reaches zero, a naive Euler–Maruyama discretisation (`schemes.tex`) can still fail: the update

$$\hat{X}_{t+\Delta t} = \hat{X}_t + \kappa(\theta - \hat{X}_t) \Delta t + \sigma \sqrt{\hat{X}_t} \sqrt{\Delta t} Z$$

has no mechanism preventing $\hat{X}_{t+\Delta t} < 0$ for an unlucky draw of Z — and once that happens, $\sqrt{\hat{X}_{t+\Delta t}}$ at the next step is not a real number at all. This is a discretisation artefact: it can occur even when the Feller condition guarantees the true continuous-time process stays positive, simply because a discrete-time approximation is not required to inherit that guarantee.

Definition 5.2 (Full truncation scheme). Write $x^+ = \max(x, 0)$. The full truncation Euler scheme `?` replaces every appearance of $\hat{X}_t$ inside the drift and diffusion evaluations — but not the running value being updated itself — by its positive part:

$$\hat{X}_{t+\Delta t} = \hat{X}_t + \kappa(\theta - \hat{X}_t^+) \Delta t + \sigma \sqrt{\hat{X}_t^+} \sqrt{\Delta t} Z.$$

Remark 5.8 (Why both terms are truncated, not just the diffusion). An earlier, more obvious fix (partial truncation) replaces $\hat{X}_t$ by $\hat{X}_t^+$ only inside the square root, leaving the drift’s mean-reversion term $\kappa(\theta - \hat{X}_t)$ using the raw (possibly negative) value. `?` found this still carries a systematic bias, and their contribution was to show that truncating the drift’s evaluation as well — `??` — reduces that bias further, at no extra implementation cost. The running value $\hat{X}_t$ itself is still allowed to go negative and stay negative for a stretch (it is only clipped where it is used, inside the coefficient evaluations); this is a deliberate feature of the scheme, not an oversight, and is what distinguishes “full” from “partial” truncation.

Remark 5.9 (Correspondence with the code). This is the discretisation `hasklib`’s roadmap specifies for CIR (OU $\rightarrow$ ABM $\rightarrow$ CIR with full truncation $\rightarrow$ Milstein), in preference to plain `EulerMaruyama::step` applied naively: wherever `CIR::drift`/`CIR::diffusion` are evaluated inside a stepping scheme, the state passed in should be $\hat{X}_t^+$, not $\hat{X}_t$ directly. Establishing the strong-order convergence properties of this scheme is a separate question from motivating it and is not addressed here.

Chapter 6

Numerical Schemes: Euler–Maruyama and Milstein